

## Problem Set 4 for lecture Mining Massive Datasets

Due November 18, 2024, 23:59 CET

---

### Exercise 1

(2 points)

Consider a cluster of  $n$  machines. Each has a probability  $p$  of failing in a given period of time  $T$ .

- a) What is the probability of at least one machine failure during this period of time?

**Hint:** Consider the complementary event that no machine fails during that time.

- b) For  $0 \leq k \leq n$ , what is the probability  $p_k$  (in terms of  $k, n$ , and  $p$ ), of exactly  $k$  machines failing during  $T$ ? Give an explanation.
- c) Show that the probabilities from Part b) satisfy:  $p_1 + p_2 + \dots + p_n = 1 - (1 - p)^n$ .

### Exercise 2

(1 point)

Give two examples for hash functions as alternatives to the arithmetic remainder function `mod`. See this video<sup>1</sup> which was created in the weeks after Adobe's password breach in October 2013 and explain the connection of hash functions to password security. What is a rainbow table? What is salt in this context?

### Exercise 3

(2 points)

Explain the details of the data structure `HashMap` from Java standard library by answering the following questions (to this aim, analyze the source code and/or documentation of the implementation).

- Describe briefly the internal data structure used by `HashMap` for storing of map entries.
- How is a new `<key, value>` pair added to a map? Consider the possibility that a collision of hash values may occur.
- How is an entry retrieved given a search key? Consider the possibility that a bucket can have more than one entry.

### Exercise 4

(1 point)

Suppose you have to design/implement hash functions for a Bloom filter algorithm. Propose 3 hash functions which are as independent from each other as possible and briefly justify your choice.

---

<sup>1</sup>Computerphile: How NOT to store passwords! <https://www.youtube.com/watch?v=8ZtInClXe1Q>

**Exercise 5****(2 points)**

Consider a Bloom filter with a bit array of  $n = 5$  bits.

- a) Compute the probability that a random element gets hashed by a hash function to a given bit in the bit array by using the probability formula from Lecture 5 (slide *Analysis: Throwing Darts (2)*). Explain your answer.
- b) Now we want to use the following hash functions ( $k = 2$ ):  
$$h_1(x) = x \bmod 5,$$
$$h_2(x) = 2 \cdot x + 3 \bmod 5.$$

Is every bit equally likely to be hit by one of the two hash functions? Suppose the numbers 4 and 1 are in the set that we want to find with the filter. Show the state of the bit array after you apply the hash functions to these numbers.
- c) Consider again the hash functions and the bit array state as in Part b). How likely is it for a number in the stream to be a false positive? Use the false positive probability formula from Lecture 5 (slide *Bloom Filter -- Analysis*). Explain your answer.

**Exercise 6****(2 points)**

Consider the Bloom filtering technique and the set of keys  $S$  with  $m$  members. Assume that you can control the size  $n$  of the bit array (bitset) used in the filter ( $n \geq m$ ) in steps  $m/8$ .

- a) Plot the optimal values of the number of hash functions  $k$  (as integers) for  $n = m, 2m, \dots, 20m$ . Describe your observations, if any. (0.5 points)
- b) Assume that you want to minimize the size  $n$  (which drives up your memory cost) for a given upper bound  $b$  on the expected false positive probability. Determine the optimal values for  $n$  (and so  $k$ ) under this objective for  $b_1 = 0.01$ ,  $b_2 = 0.005$ , and  $b_3 = 0.001$ . Also state the (expected) false positive probability for each of these parameter combinations. Explain your approach and explain why you do not need to know the actual value of  $m$ .

**Hint:** you might need to write a short program to perform the concrete calculation. (1.5 points)

**Exercise 7****(2 points)**

Consider again the Bloom filtering technique. Prove that the optimal number  $k_{opt}$  of hash functions for given size  $n$  of the bitset and a given number  $m$  of keys is  $k_{opt} = n/m \ln(2)$  (as claimed in the Lecture 5, slide *Bloom Filter -- Analysis (2)*).

**Exercise 8****(2 points)**

Suppose we have a stream of tuples with the schema `Grades(university, courseID, studentID, grade)`. Assume universities are unique, but a courseID is unique only within a university (i.e., different universities may have different courses with the same ID, e.g., "IP1") and likewise, studentIDs are unique only within a university (different universities may assign the same ID to different students). Suppose we want to answer certain queries approximately from a 1/20th sample of the data. For each of the queries below, indicate how you would construct the sample. That is, tell what the key attributes should be.

- a) For each university, estimate the average number of students in a course. (1 point)
- b) Estimate the fraction of students who have an average grade of 2.0 or better. (0.5 points)
- c) Estimate the fraction of courses where at least half of the students got the grade 1.7 or better. (0.5 points)

## Exercise 9

(2 points)

Analyze the WordCount example for Spark Structured Streaming from the lecture, also described in the Structured Streaming Programming Guide as “Quick Example”<sup>2</sup>. Run (under Linux) the sample code as described in the guide:

```
spark-submit structured_network_wordcount.py localhost 9999
```

Then, in another terminal window, run the following commands:

```
mkfifo fifo
cat > fifo &
nc -lk 9999 < fifo &
```

Now the *netcat* process is running in the background, listening on the named pipe **fifo**. All that is sent through this named pipe will be forwarded by *netcat* to port 9999, where in turn the Spark Structured Streaming program is listening. When you forward the standard output of any process with the help of the **>** operator to **fifo**, this output will be streamed to the **structured\_network\_wordcount.py** program.

- a) Run the command

```
echo "A small step for a man, a giant leap for mankind." \
> fifo
```

and describe what happens. Does the output of the streaming program behave as expected? (0.25 points)

- b) Run the command

```
yes "k thx bai bai" > fifo
```

What do you notice? How large is roughly the throughput when you assume one byte per character? (0.25 points)

- c) Modify the sample code to compute: (i) the word lengths and (ii) the average length of words received from the network in real-time. Display both the count of each word length and the average word length in the console, updated continuously as new data arrives. Name your code as **structured\_network\_wordlength.py**, test and submit it with your solution. (1.25 points)

**Hint:** You might use **length()** and **agg()** functions.

- d) Run your **structured\_network\_wordlength.py** as described above (i.e., using **fifo**) on “*The Complete Works of William Shakespeare*”<sup>3</sup>. Report your results. (0.25 points)

<sup>2</sup><https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#quick-example>

<sup>3</sup>Available here <https://www.gutenberg.org/files/100/100-0.txt> and in *heiBOX* under *MMD-Exercises\data* (consider only the content between “\*\*\* START OF THE PROJECT ... \*\*\*” and “\*\*\* END OF THE PROJECT ... \*\*\*”)